

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rc
```

```
In [2]: %matplotlib inline
from __future__ import division
rc('font', **{'family': 'sans-serif', 'sans-serif': ['Helvetica']})
rc('text', usetex=True)
```

Mid-term exam

1.

Consider the joint density $f(x, u) = 1(0 < u < \pi(x))$, the conditional density should be:

- given x , u should satisfy a uniform distribution $U(0, \pi(x))$
- given u , x should satisfy a uniform distribution on $A_u = \{x : \pi(x) > u\}$

Thus, the algorithm of the Gibbs sampler:

- sample x_0 from a proposed initial density $q_0(x)$
- sample u from $U(0, \pi(x))$
- sample x from $U(A_u)$
- iterate 2 and 3 until convergence is reached and enough samples are generated

Specifically, to sample $\pi(x) \propto 1/(1 + x^4)$

```
In [3]: def pi(x):
return 1/(1+x**4)
```

From the above algorithm and joint density, $u|_x \sim U(0, 1/(1 + x^4))$:

```
In [4]: def draw_u(x):
u = np.random.uniform()
while u >= pi(x):
u = np.random.uniform()
return u
```

And $x|_u \sim U(-\pi^{-1}(u), \pi^{-1}(u))$, where $\pi^{-1}(u)$ is defined by:

$$\pi^{-1}(u) = \left(\frac{1}{u} - 1\right)^{1/4}$$

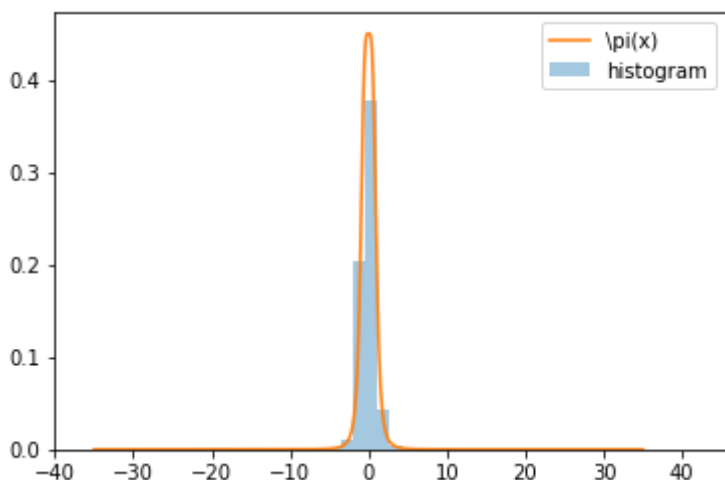
```
In [5]: def draw_x(u):
        x = np.random.uniform(-pow(1/u-1, 1/4), pow(1/u-1, 1/4))
        return x
```

Use this Gibbs sampler to generate a million samples:

```
In [6]: u = np.random.uniform()
        N = 1000000
        samples = np.zeros(N)
        for i in range(N):
            x = draw_x(u)
            samples[i] = x
            u = draw_u(x)
```

Plot the histogram of histogram of the samples, it can be shown that this is consistent with the actual distribution:

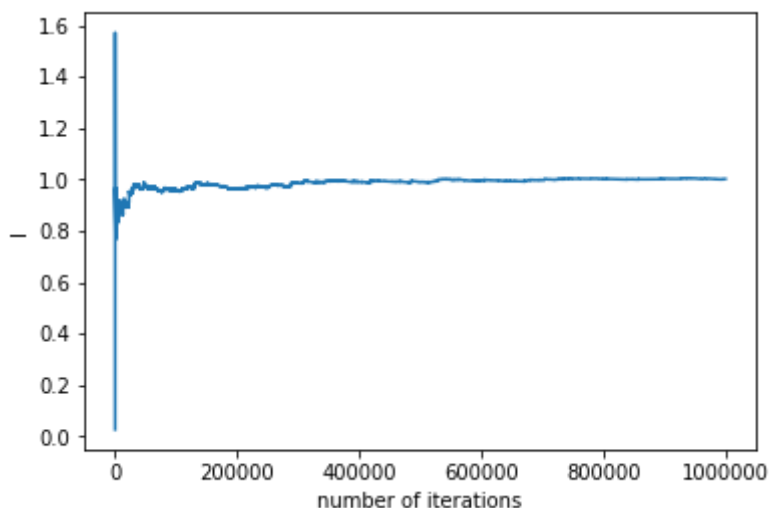
```
In [7]: sns.distplot(samples, kde=False, norm_hist=True, label="histogram")
        x = np.linspace(-35, 35, 1000)
        plt.plot(x, 1/(1+x**4) * np.sqrt(2) / np.pi, label='\pi(x)')
        plt.legend()
        plt.show()
```



Compute the Monte Carlo integral for $I = \int x^2 \pi(x) dx$, plot the convergence curve

```
In [8]: I = samples**2
        for i in range(1, N):
            I[i] = (I[i-1]*i + samples[i]**2)/(i+1) #fast, probably unstable way
            to compute average vs iteration

        plt.plot(range(N), I)
        plt.xlabel("number of iterations")
        plt.ylabel("I")
        plt.show()
```



Use the first 200000 iterations as a burn-in simulation and compute the MC approximation $\hat{I}_N = \sum_i^n x_i y_i / n$ to estimate variance:

```
In [9]: (samples**2)[200000:].mean()
```

```
Out[9]: 1.010457223164956
```

Note that I assume $\pi(x)$ is a probability function. Thus this MC integration is close to the actual integration which is:

$$I = \frac{\int x^2 \pi(x) dx}{\int \pi(x) dx} = 1$$

2.

From Bayes' theory:

$$\begin{aligned} f(\mu, \lambda \mid \text{data}) &\propto \prod_{i=1}^n g(x_i \mid \mu, \lambda) \cdot f(\mu) \cdot f(\lambda) \\ &= \lambda^{n/2} e^{\lambda \sum_{i=1}^n -(x_i - \mu)^2 / 2} \cdot \frac{1}{\tau} e^{-\mu^2 / 2\tau^2} \cdot \lambda^{a-1} e^{-\lambda b} \end{aligned}$$

Since the size of the samples is rather small, I picked large τ and small a and b so that the Markov chain is less sensitive to the coefficients:

```
In [10]: tau = 10; a = 2; b = 5
         n = 51
```

Rearrange $f(\mu, \lambda | data)$ and get conditional densities for μ and λ :

$$p(\mu | \lambda, data) \propto e^{\lambda \sum_{i=1}^n -(x_i - \mu)^2 / 2} \cdot e^{-\mu^2 / 2\tau^2}$$

$$\propto e^{-\frac{1}{2}(\mu - \frac{\lambda n \bar{x}}{1/\tau^2 + n\lambda})^2 \cdot (1/\tau^2 + n\lambda)}$$

$$\mu | \lambda \sim N\left(\frac{\lambda n \bar{x}}{1/\tau^2 + n\lambda}, \frac{1}{1/\tau^2 + n\lambda}\right)$$

```
In [11]: def draw_mu(lamda):
         return np.random.normal(lamda*10.2/(1/tau**2+n*lamda), np.sqrt(1/(n*
         lamda+1/tau**2)))
```

$$p(\lambda | \mu, data) \propto \lambda^{n/2+a-1} e^{-\frac{1}{2}\lambda[b + \sum_i (x_i - \mu)^2]}$$

$$\lambda | \mu \sim \text{Gamma}(n/2 + a, \frac{b + \sum_i (x_i - \mu)^2}{2})$$

```
In [12]: def draw_lamda(mu):
         return np.random.gamma(shape = a+n/2, scale=2/(b+(39.6 -2*mu*10.2 +
         n*mu**2) - 10.2**2/n))
```

Similar to the previous question, the Gibbs sampler starts with an initial μ_0 , then sample λ and μ iteratively based on the above two conditional densities in a round Robin fashion (100000 samples were generated):

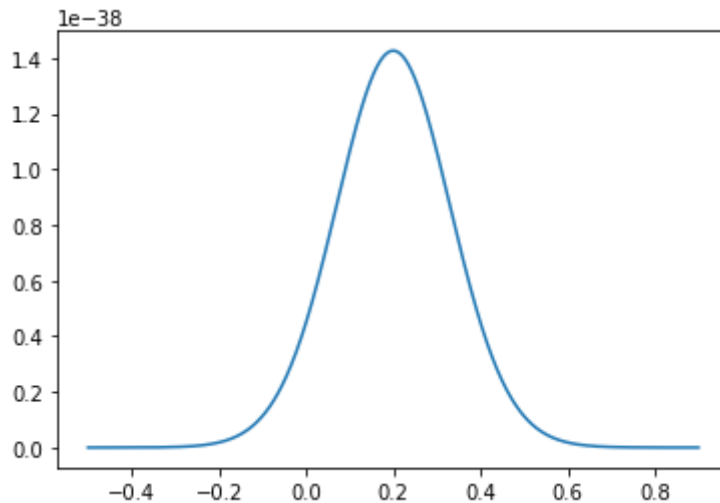
```
In [13]: mu = np.random.normal(0.2, 0.751)
         N2 = 100000
         samples2 = np.zeros((N2, 2))
         for i in range(N2):
             lamda = draw_lamda(mu)
             mu = draw_mu(lamda)
             samples2[i][0] = mu
             samples2[i][1] = lamda
```

Consider the marginal densities for μ :

$$p(\mu | data) \propto \frac{f(\mu, \lambda | data)}{p(\lambda | \mu, data)}$$

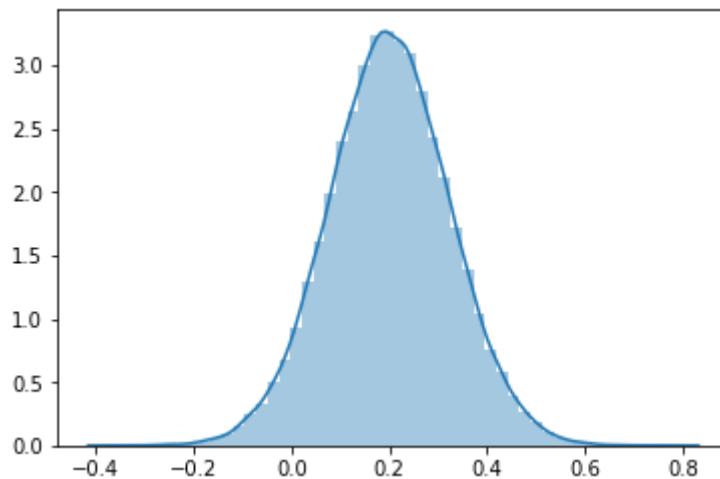
$$\propto \exp(-\mu^2 / 2\tau^2) \cdot [b + \frac{\sum_i (x_i - \mu)^2}{2}]^{-(a+n/2)}$$

```
In [14]: x = np.linspace(-0.5, 0.9, 1001)
y = np.exp(-x**2/2/tau**2)*(b + (39.6 - 2*10.2*x + n*x**2)/2)**(-a-n/2)
plt.plot(x,y, label = r"p(\mu)")
plt.show()
```



The sample distribution for μ :

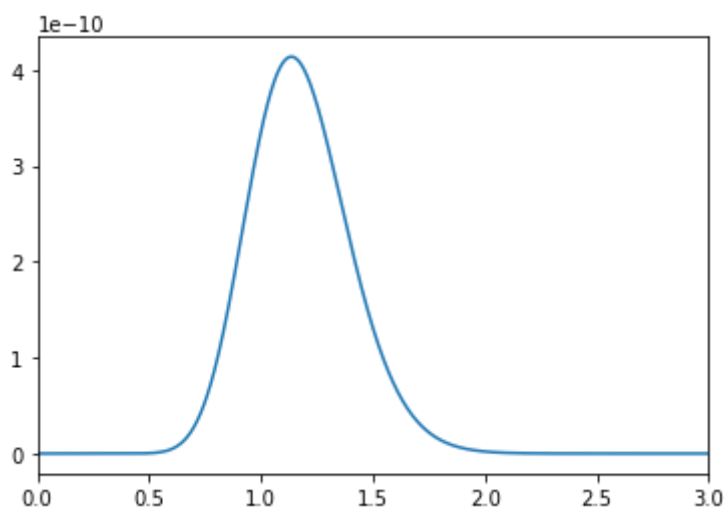
```
In [15]: sns.distplot(samples2[:,0], norm_hist=True, label="histogram")
plt.show()
```



The marginal density of λ :

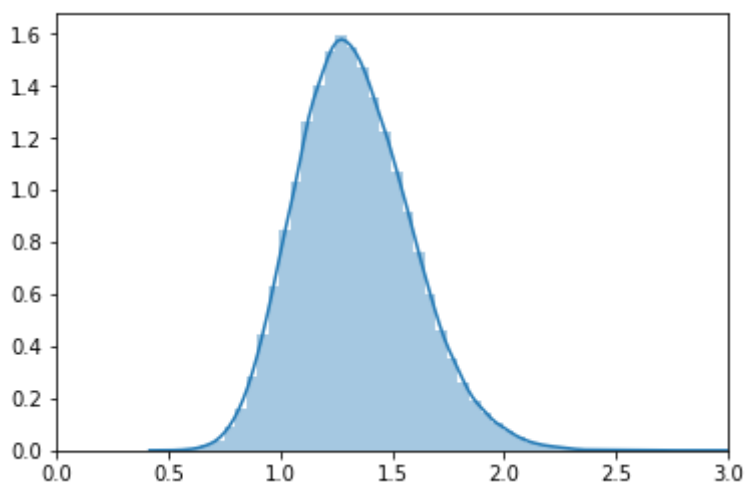
$$\begin{aligned}
 p(\lambda|data) &\propto \frac{f(\mu, \lambda | data)}{p(\mu | \lambda, data)} \\
 &\propto \lambda^{n/2+a-1} e^{-\lambda b} \sqrt{1/\tau^2 + n\lambda} e^{-\frac{1}{2}[\lambda \sum_{i=1}^n x_i^2 - \frac{(\sum_i x_i \lambda)^2}{n\lambda + 1/\tau^2}]}
 \end{aligned}$$

```
In [16]: x = np.linspace(0,3,1001)
y = x**(n/2+a-1)*np.exp(-x*b) * np.sqrt(1/tau**2+n*x) * np.exp((10.2*x)*
*2/(n*x+1/tau**2))/2 - x*39.6/2)
plt.plot(x, y)
plt.xlim(0,3)
plt.show()
```



The sample distribution of λ :

```
In [17]: sns.distplot(samples2[:,1])
plt.xlim(0,3)
plt.show()
```



The sample distribution is consistent with the real marginal distribution.

With this sampler, we can estimate the value of μ and λ by taking the sample mean:

$$\bar{\mu} = \frac{1}{N} \sum_i^N \mu_i$$

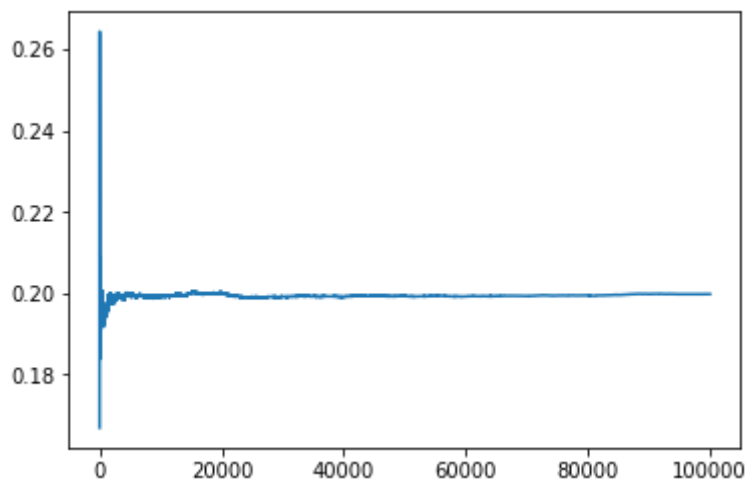
$$\bar{\lambda} = \frac{1}{N} \sum_i^N \lambda_i$$

Now plot the convergence curve of μ and λ respectively:

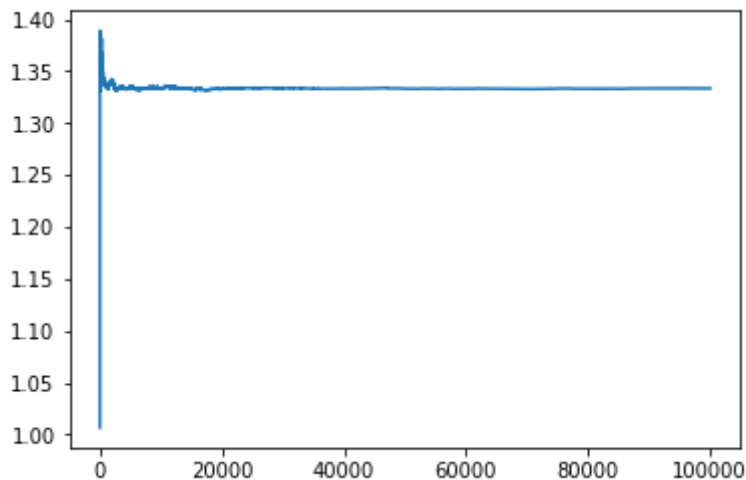
```
In [18]: mu = samples2[:,0].copy()
lamda = samples2[:,1].copy()

for i in range(1, N2):
    mu[i] = (mu[i-1]*i + samples2[i, 0])/(i+1)
    lamda[i] = (lamda[i-1]*i + samples2[i,1])/(i+1)

plt.plot(range(N2), mu)
plt.show()
```



```
In [19]: plt.plot(range(N2), lamda)
plt.show()
```



```
In [20]: samples2[-20000:].mean(axis=0)
```

```
Out[20]: array([ 0.20111011,  1.33477983])
```

$\bar{\mu}$ and $\bar{\lambda}$ from sampling the posterior with the Gibbs sampler are 0.2012 and 1.329 respectively. The mean of the data is $10.2/51 = 0.2$ and inverse variance of the data is

$1/\frac{1}{n-1} \sum_i (x_i - \bar{x})^2 = 1/[(39.6 - 2 \cdot 0.2 \cdot 10.2 + n \cdot 0.2^2)/(n-1)] = 1.3312$. Thus the estimates agree with the sample mean and variance.

3.

```
In [21]: alpha3 = 3
```

Given the power law distribution: $P(X = k) \propto k^{-\alpha}$, $k \in \{1, 2, 3, \dots\}$, $\alpha > 1$

```
In [22]: def pi(k):
          return 1/np.power(k, alpha3)
```

Proposal for Metropolis-Hasting algorithm (1D random walk):

```
In [23]: def q(x2, x1):
          assert abs(x2-x1) == 1 #check is the two states are next to each other
          if x2 == 2 and x1 == 1:
              return 1
          return 1/2
```


The $\alpha(x_2, x_1)$ for Metropolis-Hasting:

```
In [24]: def alpha(x2, x1):
          return min(1, pi(x2) / pi(x1) * q(x1, x2) / q(x2, x1))
```

Algorithm for Metropolis-Hasting:

1. Start from a x_1
2. sample x_2^* from proposal probability
3. sample u from $U(0,1)$
4. accept x_2^* if $u < \alpha(x_2^*, x_1)$, else stay at x_1
5. iterate through 2-4 until enough samples are acquired.

For this problem, we start the Markov chain at 1 and iterate 10000 times.

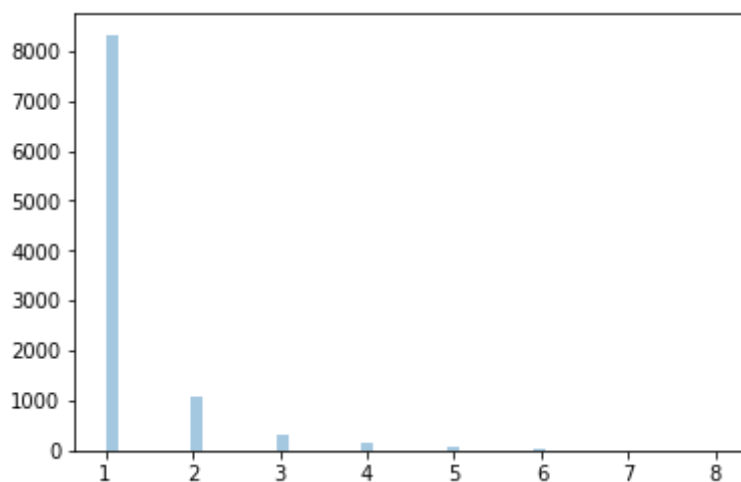
```
In [25]: x1 = 1
          N3 = 10000
          samples3 = np.zeros(N3)
          for i in range(N3):
              x2_tmp = 2 if x1 == 1 else np.random.choice([x1+1, x1-1])
              u = np.random.uniform()
              x2 = x2_tmp if u < alpha(x2_tmp, x1) else x1
              x1 = x2
              samples3[i] = x2
```

Note that the initial position matters in this problem. If the initial position x_0 is very large, $q(x_0)$ will be very close to $q(x_0 - 1)$ and $q(x_0 + 1)$, thus $\alpha(x_0 + 1, x_0)$ is also close to 1 ($\alpha(x_0 - 1, x_0) = 1$). This means that the Markov chain could walk to the right easily and takes very long iterations to converge to the true density.

Sample distribution:

```
In [26]: sns.distplot(samples3, kde = False)
```

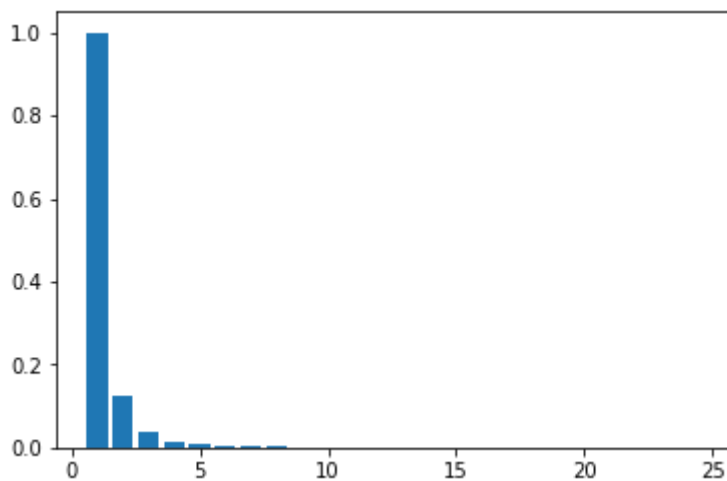
```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1d7d59b0>
```



True density function:

```
In [27]: X3 = range(1, 25)
plt.bar(X3, 1/np.power(X3, alpha3))
```

```
Out[27]: <Container object of 24 artists>
```



4.

Consider the joint density $\pi(x, y)$ and a Gibbs sampling framework is used to sample a Markov chain with stationary density π . The conditionals are $\pi(x|y)$ and $\pi(y|x)$ and while it is easy to sample $\pi(x|y)$ it is not possible to sample $\pi(y|x)$ directly.

How to proceed using a Metropolis-Hasting step:

1. start from an initial x_0, y_0
2. sample x_1 from $\pi(x|y)$ directly
3. sample y_1 from $\pi(y|x)$ using one step of Metropolis-Hasting:

i. sample y^* from proposal density $q(y^*|y)$.

ii. sample u from $U(0,1)$

iii. accept y^* if $u < \alpha(y^*, y_0)$, else stay at y_0 . In this case:

$$\alpha(y_2, y_1) = \frac{\pi(y_2|x_1) q(y_1|y_2)}{\pi(y_1|x_1) q(y_2|y_1)}$$

4. repeat 2 and 3 until enough samples are acquired.

Why does this work:

The transition density: $P(x', y' | x, y) = \pi(x' | y') \cdot P_M(y' | x, y)$

Then, consider $\pi(x', y')$:

$$\begin{aligned} \int \int P(x', y' | x, y) \pi(x, y) dx dy &= \int \int \pi(x' | y') \cdot P_M(y' | x, y) \pi(x, y) dx dy \\ &= \pi(x' | y') \int \int P_M(y' | x, y) \pi(x, y) dx dy \end{aligned}$$

Note that $P_M(y' | x, y) \pi(y | x) = P_M(y | x, y') \pi(y' | x)$

$$\begin{aligned} &= \pi(x' | y') \int \int P_M(y | x, y') \pi(y' | x) / \pi(y | x) \cdot \pi(y, x) dx dy \\ &= \pi(x' | y') \int \int P_M(y | x, y') \pi(y' | x) \pi(x) dx dy \\ &= \pi(x' | y') \pi(y') \\ &= \pi(x', y') \end{aligned}$$

Given $\pi(x, y) \propto e^{-4x} y^3 e^{-3ye^x}$. $\pi(y|x)$ is easy to sample. $y|_x \sim \text{Gamma}(4, 3e^x)$. However, $\pi(x|y) \propto e^{-(4x+3ye^x)}$ would require metropolis-hasting step. Use a lognormal proposal density to do this Metropolis step:

$$q(x_2 | x_1) = \frac{1}{x_2} e^{-(\log x_2 - \log x_1)^2 / 2}$$

```
In [28]: def pi4(x, y):
          return np.exp(-4*x-3*y*np.exp(x))

def q4(x2, x1):
    return 1/x2 * np.exp(-1/2 * (np.log(x2) - np.log(x1))**2)

def alpha4(x2, x1, y):
    return min(1, pi4(x2, y) / pi4(x1, y) * x2/x1)
```

```
In [29]: def MH_x(x, y):
          x_tmp = np.random.lognormal(np.log(x), 1)
          u = np.random.uniform()
          x = x_tmp if u < alpha4(x_tmp, x, y) else x
          return x

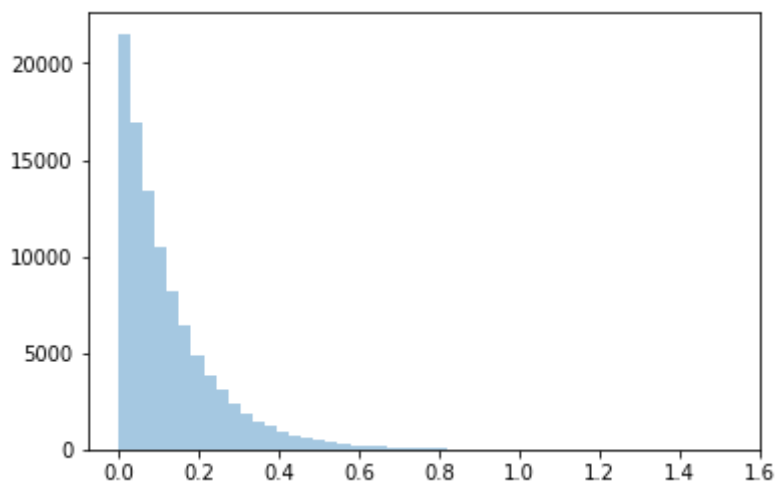
def draw_y(x):
    return np.random.gamma(shape = 4, scale = 1/3/np.exp(x))
```

```
In [30]: N4 = 100000
samples4 = np.zeros((N4, 2))
x = 1
y = np.random.gamma(4, 5)
for i in range(N4):
    y = draw_y(x)
    x = MH_x(x, y)
    samples4[i, 0] = x
    samples4[i, 1] = y
```

After 100000 iterations, the sample distribution of x:

```
In [31]: sns.distplot(samples4[:,0], kde=False)
```

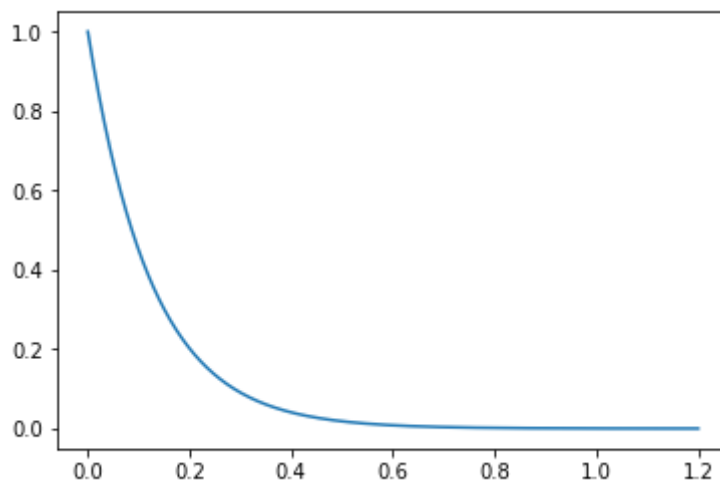
```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x1ald953898>
```



The true marginal density of $x \propto e^{-8x}$

```
In [32]: x = np.linspace(0,1.2, 1001)
plt.plot(x, np.exp(-8*x))
```

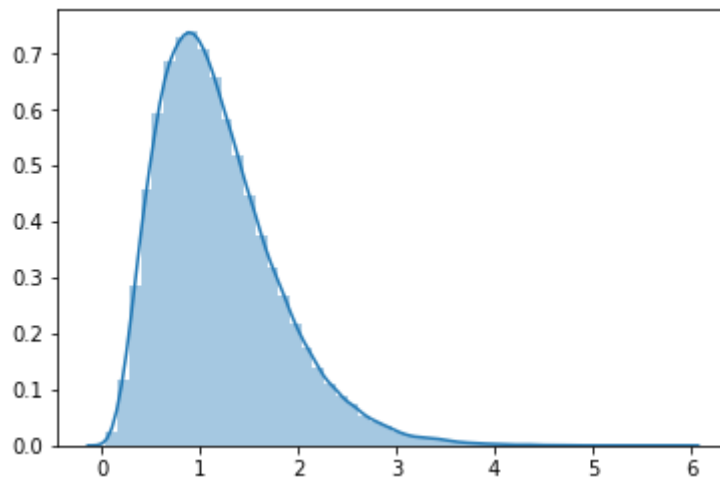
```
Out[32]: [ <matplotlib.lines.Line2D at 0x1a1da4f7f0>]
```



The sample distribution of y:

```
In [33]: sns.distplot(samples4[:, 1])
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1a1d980d30>
```

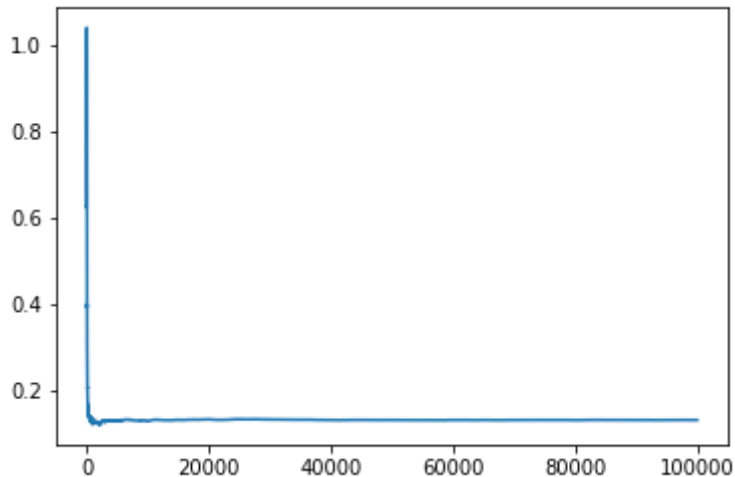


Plot the convergence of the Monte Carlo integration:

```
In [34]: I4 = (samples4[:,0] * samples4[:,1])
for i in range(1, N4):
    I4[i] = (samples4[:i,0]*samples4[:i,1]).mean()

plt.plot(range(N4), I4)
```

Out[34]: [



Use the first 20000 steps as burn-in and compute the estimated integration:

```
In [35]: print("Estimated integration: ",(samples4[20000:,0] * samples4[20000:,1]
).mean())

Estimated integration: 0.131891629783
```

This estimation also agrees with the true integrationw which is approximately 0.1317